

ZOMATO RESTAURANT INTELLIGENCE

A Complete, Nine-Phase Data Analytics & Deployment Project

9

Project Phases

41,410

Restaurants Analyzed

R^2 0.82

ML Model Accuracy

LIVE

Deployed Web Application

Yash Chavan | B.Tech CSE-DS, VIT Pune | Batch 2024-28

Live App · zomato-growth-navigator.streamlit.app | **GitHub** · github.com/yashtc2239-ops/zomato-restaurant-intelligence

01

OVERVIEW

Executive Summary

This report documents a complete, nine-phase data analytics project analyzing **41,410+ Zomato restaurants** across Bengaluru, supplemented by a 6,513-restaurant multi-city Indian subset from a second global dataset. The project moves from raw, uncleaned CSV files through business framing, ETL, exploratory analysis, SQL querying, Power BI dashboarding, machine learning, and ends in a **fully deployed, publicly accessible Streamlit web application**. Every phase was driven by a specific business question rather than generic exploration — the goal throughout was decision-grade analysis, not a classroom exercise.

Project Roadmap

#	Phase	Key Output
1	Business Understanding	Problem statement, 7 MECE business questions
2	Dataset Understanding	Schema audit of 2 Kaggle datasets (61K+ rows)
3	Data Cleaning (ETL)	41,410 + 6,513 production-clean restaurant rows
4	Exploratory Data Analysis	6 business questions answered statistically
5	Visualization	5 business-driven, insight-labeled charts
6	SQL Analysis	8 business queries — 2 linked SQL tables
7	Power BI Dashboard	3-page dashboard, 5 custom DAX measures
8	Advanced Business Analysis	Expansion Score algorithm + Random Forest ML
9	Deployment & Engineering	Streamlit + Docker + live Cloud deployment

BUSINESS & DATASET UNDERSTANDING

1.1

PHASE 1 Business Understanding

Zomato was framed as a three-sided marketplace connecting restaurants (supply), customers (demand), and the platform itself. Rather than producing generic restaurant statistics, the project was scoped around a single, decision-oriented business problem:

"Using Zomato restaurant data, identify the key drivers of customer ratings and build a pricing/performance benchmark that helps the business development team prioritize restaurant partnerships in underserved segments."

Seven business questions were defined across four MECE themes:

Theme	Business Questions
Demand & Customer Behavior	Cuisine vs order volume by city tier; price vs rating correlation
Supply & Restaurant Performance	Restaurant density by city; attributes that predict higher ratings
Operational Efficiency	Delivery time by city/cuisine/time; online order + booking combo
Strategic / BI	Which cities show high demand but low supply — where to expand

1.2

PHASE 2 Dataset Understanding

Two public Kaggle datasets were audited and combined to give both **depth** (city-level granularity) and **breadth** (multi-city comparison) — something neither dataset offered alone.

Dataset	Rows	Cols	Role in Project
Zomato Bengaluru (himanshupoddar)	51,717	17	PRIMARY — used in every phase: cleaning, EDA, SQL, Power BI, ML model, and the deployment
Zomato Global (shrutimehta)	9,551	21	SECONDARY — filtered to India (6,513 rows), used in Phase 2 schema audit, Phase 4 EDA

Audit findings: zero duplicate rows in either dataset. Bengaluru's *rate* column was stored as a string fraction (e.g. "4.1/5") with inconsistent spacing and placeholder values ("NEW", "-"). The *cost_for_two* column used comma-separated thousands ("1,200"). *dish_liked* was 54.3% null and was dropped. The Global dataset was 90.5% India records — filtering to India removed currency inconsistency with minimal data loss.

DATA CLEANING & EXPLORATORY ANALYSIS

1.3

PHASE 3 Data Cleaning (ETL)

Issue	Decision	Reason
rate stored as "4.1/5"	Regex extract to float	Enables numeric aggregation
dish_liked 54% null	Drop column	Too sparse to be reliable
cost has commas ("1,200")	Strip + cast to float	Required for price-bracket analysis
online_order/book_table Yes/No	Map to 1/0	Needed for correlation & ML model
Global: 90.5% India rows	Filter Country Code == 1	Removes currency inconsistency
Useless columns (url, phone, etc.)	Dropped	Zero analytical value

Result: Bengaluru reduced from 51,717 → 41,410 rows (17 → 11 columns). Global (India) reduced from 9,551 → 6,513 rows (21 → 14 columns), standardized to match the Bengaluru column naming convention for direct comparison.

1.4

PHASE 4 Exploratory Data Analysis

Six business questions were answered statistically on the cleaned Bengaluru dataset:

Question	Finding
Rating distribution	43.6% in Good (3.5-4.0) tier; only 1.4% reach Excellent (4.5+)
Online order vs rating	Negligible — 3.72 vs 3.66 avg (0.06 point difference)
Cost vs rating	Pearson $r = 0.38$ (moderate) — Luxury 4.16 vs Budget 3.57
Location vs rating	Lavelle Road leads at 4.14; Bommanahalli trails at 3.19
Cuisine popularity vs quality	North Indian dominates supply but ranks outside top 10 by rating
Table booking vs rating	Strongest predictor — 4.14 vs 3.62 avg, 2.6x higher cost

★ **Cross-dataset insight:** the Global (India) subset showed a lower average rating (3.35 vs 3.70) and higher average cost (■716 vs ■603) — confirming Bengaluru as a relatively higher-rated, lower-cost market within India's broader restaurant landscape.

VISUALIZATION & SQL ANALYSIS

1.5

PHASE 5 Visualization

Five matplotlib/seaborn charts were produced, each answering one explicit business question rather than being decorative. Every chart carries a one-line business insight as a subtitle — the "so what?" principle applied throughout.

Chart File	Business Question Answered
chart1_rating_distribution.png	How are restaurants distributed across rating tiers?
chart2_table_booking.png	Does table booking indicate a premium segment?
chart3_cost_vs_rating.png	Does spending more guarantee a better experience?
chart4_location_rating.png	Which locations should Zomato prioritize or improve?
chart5_cuisine_analysis.png	Which cuisines are oversupplied vs underserved?

1.6

PHASE 6 SQL Analysis

Both cleaned datasets were loaded into SQLite as two tables — *bengaluru* and *global_india* — and queried with 8 production-style business queries (also available as MySQL-compatible syntax).

Query	Business Use
Q1 — Restaurant type distribution	Understand Zomato's product/listing mix
Q2 — Top 10 locations by rating	Target areas for premium brand partnerships
Q3 — Online order impact on rating	Confirms online ordering alone doesn't drive quality
Q4 — Table booking premium analysis	Identify the premium restaurant tier
Q5 — Cuisine performance	Surface underserved high-quality cuisine segments
Q6 — Demand/supply ratio by location	Identify expansion-ready vs saturated markets
Q7 — Restaurant type performance	Which restaurant formats perform best
Q8 — Expansion recommendation	Final ranked list of HIGH/MEDIUM/SATURATED zones

★ **Hero finding (Q6/Q8): Church Street and Lavelle Road show 1,000+ votes per restaurant — more than double the city average — flagging both as demand-starved, expansion-ready zones.**

POWER BI DASHBOARD

A 3-page interactive Power BI dashboard was built on the cleaned Bengaluru dataset, using 5 custom DAX measures to move beyond simple SUM/COUNT aggregations into comparative and relative-performance metrics — the kind of work that separates a tutorial project from an industry-level deliverable.

Dashboard Pages

Page	Contents
1 · Executive Summary	4 KPI cards, rating distribution, order/booking donuts, cost-bracket slicer
2 · Location Intelligence	Top/Bottom 10 locations, cost-vs-rating scatter, Expansion Priority table
3 · Cuisine & Business	Supply vs Quality comparison, table-booking cards, Value Score rankings

DAX Measures

Measure	Purpose
City Avg Rating	Context-independent baseline using ALL() for fair comparisons
Rating Variance vs City	Shows how far each location deviates from the city baseline
Value Score	Price-tier-adjusted rating using ALLEXCEPT — surfaces hidden gems
Performance Flag	Conditional Top Performer / Underperformer tag
Demand Pressure Index	Votes-per-restaurant relative to city average — the expansion signal

DAX measures using **ALL()** and **ALLEXCEPT()** filter-context manipulation are a common differentiator between tutorial-level and industry-level Power BI work — demonstrating fluency beyond drag-and-drop visuals.

PHASE 08

ADVANCED BUSINESS ANALYSIS

Expansion Score Algorithm

Expansion Score = (40% × Demand Pressure) + (35% × Avg Rating) + (25% × Supply Gap)

All three metrics normalized 0–100 via MinMaxScaler across 72 locations with 30+ restaurants. **3 zones classified Critical priority** (score > 80).

Rank	Location	Score	Priority	Rest.	Rating
1	Lavelle Road	95.6	CRITICAL	481	4.14
2	Church Street	91.1	CRITICAL	546	3.99
3	St. Marks Road	81.6	CRITICAL	343	4.02
4	Koramangala 3rd Block	78.1	HIGH	191	4.02
5	Koramangala 4th Block	76.2	HIGH	841	3.92

Rating Prediction Model

Model	R² Score	RMSE	Verdict
Linear Regression	0.2920	0.3696	Weak — linear assumption insufficient
Random Forest	0.8157	0.1885	WINNER — 179% better R²

Feature	Importance	Interpretation
votes (engagement)	69.3%	DOMINANT — 4.7x more important than price
cost_for_two	14.7%	Moderate influence on quality
rest_type	10.9%	Format matters — fine dining rates highest
online_order	3.0%	Minimal direct impact
book_table	2.1%	Confirms premium segment signal

★ Key Finding: Customer engagement (votes) is 4.7x more important than price in predicting restaurant quality — Zomato's ranking algorithm should weight engagement signals higher than price-based factors.

DEPLOYMENT & ENGINEERING

The final phase converts the analysis into a production-style deployed product rather than leaving it as static notebooks — the difference between a data analyst portfolio and an engineer's portfolio.

Streamlit Web Application — Live & Public

Page	Features
Executive Overview	4 live KPI cards, rating distribution, service donuts, cost-vs-quality chart
Location Intelligence	Top/Bottom 10 locations, color-coded Expansion Priority table (live-filtered)
Cuisine & Business	Supply vs Quality charts, table-booking premium signal comparison
ML Model Insights	Feature importance, model comparison table, Live Rating Predictor widget

Sidebar filters (Location, Cuisine, Price Range, Min Rating) dynamically re-filter all four pages in real time, including re-computation of the Expansion Score and ML feature charts. Edge cases — such as a filter combination returning zero or near-zero restaurants — are handled gracefully with informative messages rather than runtime errors.

Containerization (Docker)

Layer	Purpose
FROM python:3.11-slim	Pinned, lightweight base image for reproducibility
WORKDIR /app	Isolated, clean working directory inside the container
COPY requirements.txt + RUN pip install	Dependency layer cached separately from app code
COPY . .	Application code layer — changes most frequently
EXPOSE 8501	Documents the Streamlit port for cloud routing
HEALTHCHECK	Lets cloud platforms detect a crashed container
CMD streamlit run ...	Entry point — binds to 0.0.0.0 for external access

Layer ordering follows Docker best practice: dependencies are installed before application code is copied in, so code edits do not invalidate the (slow) pip install cache layer.

Cloud Deployment — Continuous Delivery

The application is deployed on Streamlit Community Cloud, connected directly to the GitHub repository. Every git push to the main branch triggers an automatic rebuild and redeploy — a lightweight CI/CD pipeline with zero manual deployment steps after initial setup.

Item	Detail
Live URL	zomato-growth-navigator.streamlit.app
Source repo	github.com/yashtc2239-ops/zomato-restaurant-intelligence
Deployment model	Streamlit Cloud — auto-redeploy on every git push

Local alternative	Docker — docker build + docker run, identical to cloud behavior
-------------------	---

BUSINESS RECOMMENDATIONS & DELIVERABLES

1

EXPAND in 3 Critical Priority Zones

Lavelle Road (95.6), Church Street (91.1), St. Marks Road (81.6) — high demand pressure, strong ratings, relatively low supply. Onboard new restaurant partners here before saturation.

2

REDESIGN the Rating Algorithm

Customer votes drive 69% of rating prediction — 4.7x more than cost. Reweight engagement signals in the ranking algorithm to improve recommendation quality.

3

TARGET the Premium Cuisine Gap

Modern Indian (4.31), European (4.26), Mediterranean (4.20) outperform in ratings but have low restaurant counts — high demand, low competition opportunity.

4

SCALE the Table Booking Program

Table booking restaurants rate 0.52 points higher and charge 2.6x more (■ 1,276 vs ■ 482). Expanding Zomato Gold/Pro captures meaningful revenue upside.

5

QUALITY IMPROVEMENT for Bottom 10 Locations

Bommanahalli (3.19), RT Nagar (3.46) have high restaurant counts but poor quality scores. Targeted training programs would lift overall platform rating.

02

DELIVERABLES Project Outputs

Deliverable	Description
Live Streamlit App	4-page interactive dashboard, deployed on Streamlit Cloud
GitHub Repository	Full source code, notebooks, SQL, reports, Docker config
Power BI Dashboard (.pbix)	3-page dashboard with 5 custom DAX measures
6 Jupyter Notebooks	Phases 1-8 fully documented, reproducible analysis
SQL Business Queries	8 queries covering Bengaluru + Global India tables
Dockerfile	7-layer container definition, production-ready
5 Chart PNGs + 2 ML charts	Expansion score, feature importance, and 5 EDA visualizations
This Report (PDF)	Full 9-phase business and technical analysis